

Chp 4 Linear Regression

Chp 5 Logistic regression

Hands-on Machine Learning with R
Booookclub R-Ladies Utrecht and R-Ladies Den Bosch

Martine Jansen

- Organized by @RLadiesUtrecht and @RLadiesDenBosch
- Meet-ups every 2 weeks on “Hands-On Machine Learning with R” by Bradley Boehmke and Brandon Greenwell
- No session recording! But we will publish the slides and notes!
- We use HackMD for making shared notes and for the registry:
<https://hackmd.io/rGu7xw2bRS-lm8lq7-wvXw>
- Please keep mic off during presentation. Nice to have camera on and participate to make the meeting more interactive.
- Questions? Raise hand / write question in HackMD or in chat
- Remember presenters are not necessarily also subject experts 😊
- Remember the R-Ladies code of conduct.
In summary, please be nice to each other and help us make an **inclusive** meeting! ❤️

What did we talk about last time?



- Target engineering: transform outcome variable via `log/Box-cox`
- Missingness: informative/random, imputation via estimated statistic/KNN
- Feature filtering: remove (near)-zero variance variables
- Numeric feature engineering: sometimes useful to transformation to reduce skewness, standardization
- Categorical feature engineering: lumping, one-hot / dummy encoding, label encoding
- Dimension reduction: see chp 17-19
- Proper implementation: sequential steps, data leakage, recipes

Part II Supervised Learning

Chp 4 Linear Regression

Approximate (linear) relationship between **continuous** response variable and set of predictor variables

4.1 Prerequisites

Libraries

```
1 library(dplyr)      # for data manipulation
2 library(ggplot2)    # for graphics
3
4 library(caret)       # for cross-validation, etc.
5 library(rsample)     # you have to scroll back in the book to detect
6                     # necessary for initial_split
7 library(vip)         # variable importance
8 #library(pdp)        # is used in section on variable importance
```

Code for the data, from previous chps

```
1 ames <- AmesHousing::make_ames()
2
3 set.seed(123)
4 split <- initial_split(ames, prop = 0.7,
5                       strata = "Sale_Price")
6 ames_train <- training(split)
7 ames_test  <- testing(split)
```

4.2 Simple linear regression

If Y and X are (approx) linearly related then:

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i, \text{ for } i = 1, \dots, n, \text{ and } \epsilon_i \sim N(0, \sigma^2)$$

β_0 : intercept, average response when $X = 0$

β_1 : slope, increase in average response per 1 unit increase in X


Using least squares regression, coefficients can be calculated with `lm`:

```
1 model1 <- lm(Sale_Price ~ Gr_Liv_Ar
2               data = ames_train)
3 model1$coef
```

```
(Intercept) Gr_Liv_Area
15938.1733   109.6675
```

```
1 sigma(model1)
```

```
[1] 56787.94
```



Inference

- Point estimates for β_0 , β_1 and σ not that interesting
- Need to know how much they vary
- When these assumptions are met:
 - independent obs
 - random error mean zero, constant variance
 - random error normally distributed

100(1 - α)% confidence interval: $\beta_j \pm t_{1-\alpha/2, n-p} \widehat{SE}_{\beta_j}$

```
1 confint(model1, level = 0.95)
```

```
          2.5 %      97.5 %  
(Intercept) 8384.213 23492.1336  
Gr_Liv_Area  104.920   114.4149
```

```
1 summary(model1)
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area, data = ames_train)
```

Residuals:

Min	1Q	Median	3Q	Max
-474682	-30794	-1678	23353	328183

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	15938.173	3851.853	4.138	3.65e-05 ***
Gr_Liv_Area	109.667	2.421	45.303	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 56790 on 2047 degrees of freedom

Multiple R-squared: 0.5007, Adjusted R-squared: 0.5004

F-statistic: 2052 on 1 and 2047 DF, p-value: < 2.2e-16

4.3 Multiple linear regression

- More main effects:

```
model2 <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built, data =  
ames_train)
```

- Add an interaction effect with `:`

```
model2a <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built +  
Gr_Liv_Area:Year_Built, data = ames_train)
```

- Add all the features in the set as main effects:

```
model3 <- lm(Sale_Price ~ ., data = ames_train)
```

- The analyst decides on having interaction effects in linear regression
- When interaction effect in model, have also comprising terms in model

4.4 Assessing model accuracy

For this example, “best” model: lowest RMSE via cross-validation

```
1 set.seed(123) # for reproducibility
2 (cv_model1 <- train(
3   form = Sale_Price ~ Gr_Liv_Area,
4   data = ames_train,
5   method = "lm",
6   trControl = trainControl(method = "cv", number = 10)
7 ))
```

Linear Regression

2049 samples
1 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 1843, 1844, 1844, 1844, 1844, 1844, ...

Resampling results:

RMSE	Rsquared	MAE
56644.76	0.510273	38851.99

Tuning parameter 'intercept' was held constant at a value of TRUE

The (averaged) RMSE for the 3 main effect models:

```
1 cv_model1$results$RMSE  
2 cv_model2$results$RMSE  
3 cv_model3$results$RMSE
```

```
[1] 56644.76
```

```
[1] 46865.68
```

```
[1] 41691.74
```

Interpret the cv result as:

When applied to unseen data, the predictions model 3 makes are, on average, about 41691.74 off from the actual sale price.

Looking at adjusted R^2 , as I got taught:

```
1 summary(model1)$adj.r.squared  
2 summary(model2)$adj.r.squared  
3 summary(model3)$adj.r.squared
```

```
[1] 0.5004063
```

```
[1] 0.6567094
```

```
[1] 0.9339612
```

Model 3 explains 94% of the variance in sale price

4.5 Model concerns


Be sure the assumptions hold:

- Linear relationship, if not transform
- Constant variance among residuals (homoscedasticity)
- No autocorrelation, errors are independent and uncorrelated
- More observations than predictors, if not try regularized regression
- No or little multicollinearity, if not difficult to separate out individual effects variables

4.6 Principal component regression

Address multicollinearity for instance by using Principal Components as predictors

```
1 set.seed(123)
2 cv_model_pcr <- train(
3   Sale_Price ~ .,
4   data = ames_train,
5   method = "pcr",
6   trControl = trainControl(method = "cv",
7                             number = 10)
8   preProcess = c("zv", "center", "scale"),
9   tuneLength = 100)
10
11 bestTune <- cv_model_pcr$bestTune[1,1]
12
13 ggplot(cv_model_pcr) +
14   geom_vline(xintercept = bestTune,
15             color = "red")
```




Question I have: - Why useful? It brings RMSE down, but do we get insight in importance of predictors? Different from regular regression? - I thought there are max `ncol(data)` PC's

Partial least squares

Supervised dimension reduction procedure:

- that finds new features - that not only captures most information in original features, - but **also are related to the response** - PLS places highest weight on variables most strongly related to response

```
1 set.seed(123)
2 cv_model_pls <- train(
3   Sale_Price ~ .,
4   data = ames_train,
5   method = "pls",
6   trControl = trainControl(method = "cv",
7                             number = 10)
8   preProcess = c("zv", "center", "scale",
9   tuneLength = 30
10 )
11
12 bestTune <- cv_model_pls$bestTune[1,1]
13
14 ggplot(cv_model_pls) +
15   geom_vline(xintercept = bestTune,
16             color = "red")
```



Added, how to see this best PLModel

```
1 the_best_pls <- cv_model_pls$finalModel
2 the_best_pls$coefficients
```

```
, , 1 comps
```


	.outcome
MS_SubClassOne_Story_1945_and_Older	-1148.4437622
MS_SubClassOne_Story_with_Finished_Attic_All_Ages	-147.5518230
MS_SubClassOne_and_Half_Story_Unfinished_All_Ages	-337.5525120
MS_SubClassOne_and_Half_Story_Finished_All_Ages	-865.0621444
MS_SubClassTwo_Story_1946_and_Newer	1790.3091140
MS_SubClassTwo_Story_1945_and_Older	-294.2140110
MS_SubClassTwo_and_Half_Story_All_Ages	142.9321733
MS_SubClassSplit_or_Multilevel	-116.6483406
MS_SubClassSplit_Foyer	-233.9429173
MS_SubClassDuplex_All_Styles_and_Ages	-499.7477346
MS_SubClassOne_Story_PUD_1946_and_Newer	474.0964976
MS_SubClassTwo_Story_PUD_1946_and_Newer	-529.4888895
MS_SubClassPUD_Multilevel_Split_Level_Foyer	-339.6653643
MS_SubClassTwo_Family_conversion_All_Styles_and_Ages	-463.2790399
MS_ZoningResidential_High_Density	-251.1720053

4.8 Feature interpretation

- Variable importance: identify variables most influential in model
- LR: often absolute value t-statistic for each parameter
- Difficult when having interactions and transformations
- PLS: contribution coefficients weighted proportionally to reduction RSS

Calculate VIP in PLS
(100 is most important):


```
1 vip(cv_model_pls,  
2     num_features = 20,  
3     method = "model")
```



PDP - partial dependence plots

- Plot change in average predicted value as specified feature(s) vary over their marginal distribution
- How fixed change in a predictor relates to fixed linear change in outcome, while taking into account average effect of all other features in model
- More useful in case of non-linear relationships (chp 16)

```
1 # This is NOT a ggplot!  
2 pdp::partial(cv_model_pls,  
3             "Gr_Liv_Area",  
4             grid.resolution = 20,  
5             plot = TRUE)
```





Part II Supervised Learning

Chp 5 Logistic Regression

Approximate the relationship between a **binary** response variable and a set of predictor variables

5.1 Prerequisites

Libraries


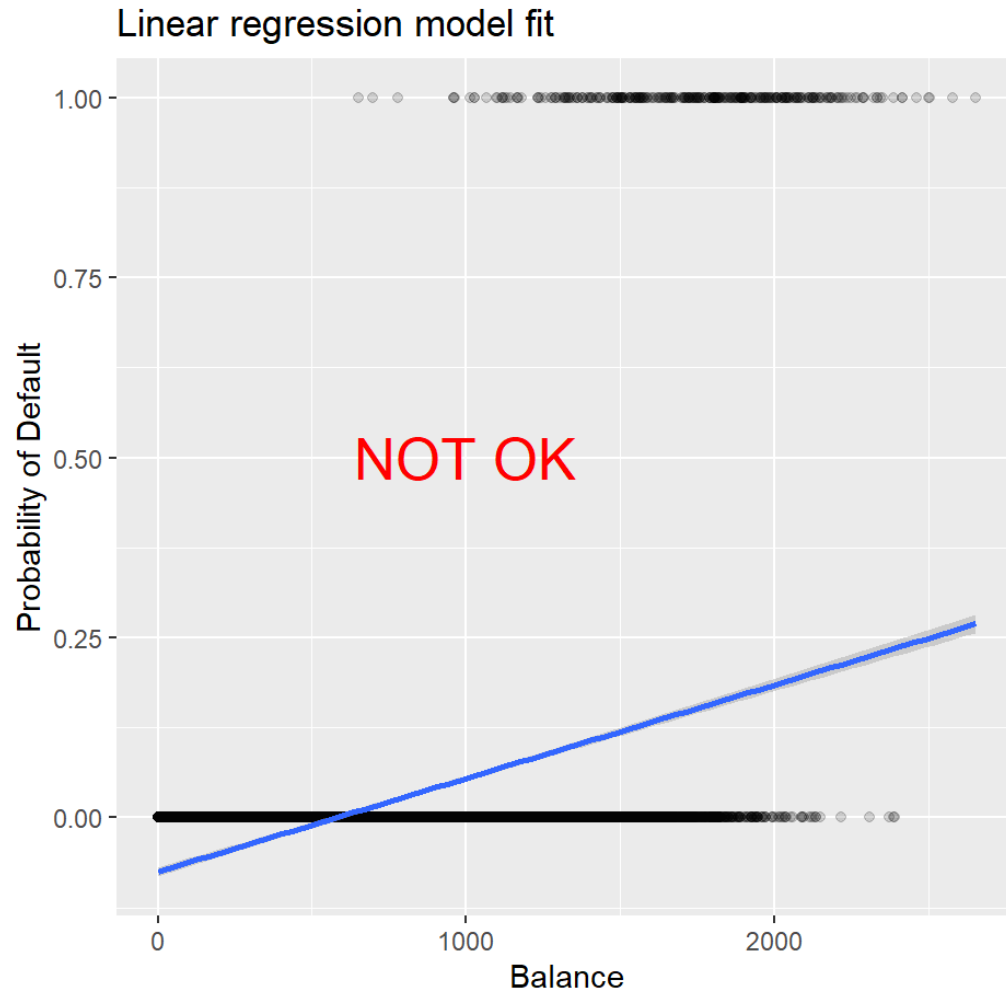
```
1 library(dplyr)      # for data manipulation
2 library(ggplot2)    # for graphics
3 library(caret)      # for cross-validation, etc.
4 library(rsample)    # necessary for initial_split
5 library(vip)        # variable importance
6 # library(modeldata)
7 # library(broom)
8 # library(ROCR)
```

Code for the data, from previous chps

```
1 # attrition <- rsample::attrition # line in book chp1 no longer works
2
3 # data are moved into the `modeldata` package
4 df <- modeldata::attrition %>%
5   # make all factors unordered
6   mutate_if(is.ordered, factor, ordered = FALSE)
7
8 set.seed(123) # for reproducibility
9 churn_split <- initial_split(df, prop = .7, strata = "Attrition")
```

R-Ladies theme for Quarto Presentations. Code available on [GitHub](#).

5.2 Why logistic regression



The formula of a sigmoid function looks complicated:

$$p(X) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}$$

Look at odds:

$$\frac{p(X)}{1 - p(X)} = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}} / \frac{1}{1 + e^{\beta_0 + \beta_1 X}} = e^{\beta_0 + \beta_1 X}$$

And then take log, and call that logit (the log of the odds):

$$\log \left(\frac{p(X)}{1 - p(X)} \right) = \log (e^{\beta_0 + \beta_1 X}) = \beta_0 + \beta_1 X$$

5.3 Simple logistic regression

Models are calculated using Maximum Likelihood


```
1 model1 <- glm(Attrition ~ MonthlyIncome,  
2               family = "binomial",  
3               data = churn_train)
```

```
1 broom::tidy(model1)[,1:2] %>%  
2   knitr::kable(booktabs = TRUE)
```

term	estimate
(Intercept)	-0.8860896
MonthlyIncome	-0.0001386

Increase of 1 unit in MonthlyIncome,

- logit of attrition 0.000139 less
- odds of attrition multiply by $\exp(-0.000139) = 0.99986$
- hence odds smaller, hence probability smaller



Confidence interval for coefficient

```
1 tidy(model1)
```

```
# A tibble: 2 x 5
  term          estimate std.error statistic    p.value
<chr>         <dbl>     <dbl>     <dbl>    <dbl>
1 (Intercept)  -0.886      0.157      -5.64 0.0000000174
2 MonthlyIncome -0.000139 0.0000272    -5.10 0.000000344
```

```
1 # for the logit coefficients:
2 confint(model1)
```

```
                2.5 %      97.5 %
(Intercept) -1.1932606571 -5.761048e-01
MonthlyIncome -0.0001948723 -8.803311e-05
```

```
1 # for the odds coefficients:
2 exp(confint(model1))
```

```
                2.5 %      97.5 %
(Intercept)  0.3032309 0.5620835
MonthlyIncome 0.9998051 0.9999120
```


5.4 Multiple logistic regression

Explaining attrition from MonthlyIncome and Overtime:

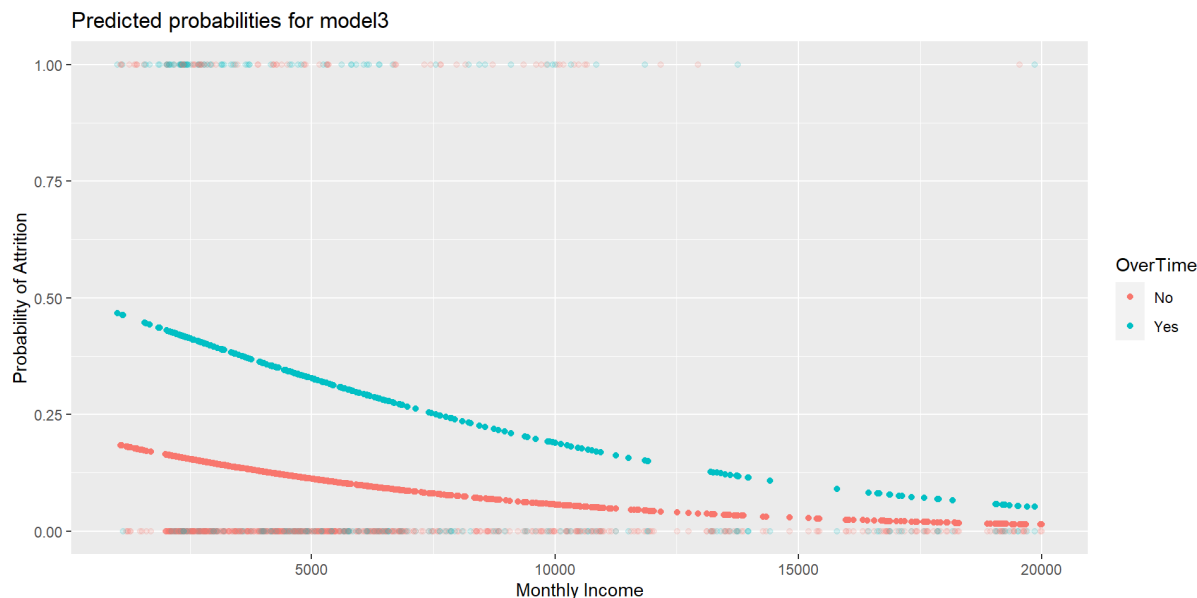
```
1 model3 <- glm(  
2   Attrition ~ MonthlyIncome + OverTime,  
3   family = "binomial",  
4   data = churn_train  
5 )  
6  
7 broom::tidy(model3)
```

```
# A tibble: 3 x 5  
  term          estimate std.error statistic  p.value  
  <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
1 (Intercept)   -1.33      0.177     -7.54 4.74e-14  
2 MonthlyIncome -0.000147  0.0000280  -5.27 1.38e- 7  
3 OverTimeYes    1.35      0.180      7.50 6.59e-14
```

```

1 churn_train3 <- # different from book:
2   # adds column "pred" to data
3   # with probs according to model 3
4   modelr::add_predictions(churn_train, model = model3, type = "response") %>%
5   mutate(prob = ifelse(Attrition == "Yes", 1, 0))
6
7 # also different from book
8 ggplot(churn_train3,
9       aes(x = MonthlyIncome, color = OverTime)) +
10  geom_point(aes(y = prob), alpha = .15) +      # observations
11  geom_point(aes(y = pred)) +                  # predictions
12  labs(title = "Predicted probabilities for model3",
13       x = "Monthly Income",
14       y = "Probability of Attrition")

```



5.5 Assessing model accuracy - how well models predict

Attrition ~ MonthlyIncome

```
1 set.seed(123)
2 cv_model1 <- train(
3   Attrition ~ MonthlyIncome,
4   data = churn_train,
5   method = "glm",
6   family = "binomial",
7   trControl = trainControl(method = "cv",
8                             number = 10)
9
10 pred_class1 <- predict(cv_model1,
11                        churn_train)
12
13 confusionMatrix(
14   data = relevel(pred_class1,
15                  ref = "Yes"),
16   reference =
17     relevel(churn_train$Attrition,
18             ref = "Yes")
19 )
```

Attrition ~ .

```
1 set.seed(123)
2 cv_model3 <- train(
3   Attrition ~ .,
4   data = churn_train,
5   method = "glm",
6   family = "binomial",
7   trControl = trainControl(method = "cv",
8                             number = 10)
9
10 pred_class3 <- predict(cv_model3,
11                        churn_train)
12
13 confusionMatrix(
14   data = relevel(pred_class3,
15                  ref = "Yes"),
16   reference =
17     relevel(churn_train$Attrition,
18             ref = "Yes")
19 )
```

Attrition ~ MonthlyIncome

Confusion Matrix and Statistics

```

Reference
Prediction Yes  No
Yes      0    0
No     165  863

Accuracy : 0.8395
95% CI : (0.8156, 0.8614)
No Information Rate : 0.8395
P-Value [Acc > NIR] : 0.5208

Kappa : 0

McNemar's Test P-Value : <2e-16

Sensitivity : 0.0000
Specificity : 1.0000
Pos Pred Value : NaN
Neg Pred Value : 0.8395
Prevalence : 0.1605
Detection Rate : 0.0000
Detection Prevalence : 0.0000
Balanced Accuracy : 0.5000

```

Attrition ~ .

Confusion Matrix and Statistics

```

Reference
Prediction Yes  No
Yes      83   20
No     82  843

Accuracy : 0.9008
95% CI : (0.8809, 0.9184)
No Information Rate : 0.8395
P-Value [Acc > NIR] : 8.982e-09

Kappa : 0.5658

McNemar's Test P-Value : 1.542e-09

Sensitivity : 0.50303
Specificity : 0.97683
Pos Pred Value : 0.80583
Neg Pred Value : 0.91135
Prevalence : 0.16051
Detection Rate : 0.08074
Detection Prevalence : 0.10019
Balanced Accuracy : 0.73993

```

No Information Rate : 0.8395: Predict most common outcome (“No”) for all, still accuracy 83.9%.

Accuracy: $P(\text{pred} = \text{actual}), (TP+TN)/(TP+FP+TN+FN)$

Sensitivity (recall): $P(\text{pred} = \text{“yes”} | \text{actual} = \text{“yes”}), TP / (TP + FN)$

Specificity: $P(\text{pred} = \text{“no”} | \text{actual} = \text{“no”}), TN / (TN + FP)$


Pos Pred Value (precision): $P(\text{actual} = \text{“yes”} | \text{pred} = \text{“yes”}), TP / (TP + FP)$

Neg Pred Value: $P(\text{actual} = \text{“no”} | \text{pred} = \text{“no”}), TN / (TN + FN)$

Prevalence: $(TP+FN)/(TP+FN+FP+FN)$

ROC curve

```
1 library(ROCR)
2
3 m1_prob <- predict(cv_model1,
4                   churn_train, type = "prob")$Yes
5 m3_prob <- predict(cv_model3,
6                   churn_train, type = "prob")$Yes
7
8 # Compute AUC metrics for models
9 perf1 <- prediction(m1_prob,
10                   churn_train$Attrition) %>%
11   performance(measure = "tpr",
12               x.measure = "fpr")
13 perf2 <- prediction(m3_prob,
14                   churn_train$Attrition) %>%
15   performance(measure = "tpr",
16               x.measure = "fpr")
17
18 plot(perf1, col = "black", lty = 2)
19 plot(perf2, add = TRUE, col = "blue")
20 legend(0.8, 0.2, legend = c("cv_model1", "cv_model3"),
21       col = c("black", "blue"), lty = 2:1, cex = 1.2)
```



Other options for ROC curves:

<https://rviews.rstudio.com/2017/05/10/r-packages-for-roc-curves/>

5.6 Model concerns

- Also important to check adequacy
- Concept of residual is difficult
- Some literature referrals

5.7 Feature interpretation



```
1 vip(cv_model3, num_features = 20)
```

5.8 Final thoughts

- Logistic regression suffers also from the many assumptions (i.e. linear relationship of the coefficient, multicollinearity)
- Often more than two classes to predict (multinomial classification)
- Future chapters discuss more advanced algorithms for binary and multinomial classification